

BIAS-VI API: Cultural Safety Middleware for AI

The ETV Institute | Center for Decision Science

License

MIT

python

3.11+

FastAPI

0.115.0

Overview

The **BIAS-VI API** is an open-source ethical safety middleware designed for Artificial Intelligence deployed in Community Violence Intervention (CVI), social services, and high-stakes urban environments.

Developed by **Barbara D. Gaskins** at the ETV Institute, this API acts as a "Cultural Safety Layer." It prevents standard LLMs from generating false positives based on street vernacular and ensures all AI responses adhere to a strict **Anti-Surveillance Guarantee**.

Core Features

- **The Translator Classifier (/v1/analyze)**: Distinguishes between culturally expressive language (grief, venting) and actionable, imminent threats.

- **The De-escalation Router (/v1/respond)**: Generates culturally competent, trauma-informed responses that prioritize the "Pause for Peace" and never recommend predictive policing.
- **The CWI Data Extraction Engine (/v1/extract)**: Extracts anonymized, macro-level data from unstructured field notes to power the Community Wellness Index (CWI) without capturing PII.

Why This Exists (The Kaggle Benchmark)

Standard LLMs (like GPT-4 and Claude) often fail when processing urban vernacular, flagging expressive grief as "violent" or recommending carceral solutions (police intervention) for community conflicts. The BIAS-VI API solves this "cultural blind spot" by applying the 5 Pillars of Intervention before the AI responds.

Quick Start

1. Clone the Repository

Bash

```
git clone https://github.com/your-org/bias-vi-api.git
cd bias-vi-api
```

2. Install Dependencies

Bash

```
pip install -r requirements.txt
```

3. Set Environment Variables

You will need an OpenAI API key (or compatible proxy) to run the core logic.

Bash

```
export OPENAI_API_KEY="your-api-key-here"
```

4. Run the Server

Bash

```
uvicorn main:app --host 0.0.0.0 --port 8000 --reload
```

The API will be available at `http://127.0.0.1:8000` . You can view the interactive Swagger documentation at `http://127.0.0.1:8000/docs` .

Testing

Run the included test suite to validate the endpoints against real-world scenarios:

Bash

```
python test_api.py
```

The Anti-Surveillance Guarantee

By utilizing this API, developers agree to the ETV Institute's Anti-Surveillance Terms of Service: **This tool must never be used to generate predictive policing targets, recommend carceral action, or share data with law enforcement without a warrant.** We track decisions, not people.

Contributing

We welcome contributions from data scientists, CVI practitioners, and engineers. Please see `CONTRIBUTING.md` for guidelines.

License

This project is licensed under the MIT License - see the `LICENSE` file for details.